

Homework 1: Multi-Layer Perceptrons and Training

Instructor: Sid Nadendla

Due: Mar 9, 2025

Goals and Directions:

- The main goal of this assignment is to implement a multi-layer perceptron from scratch, and train it on a given dataset
- Comprehend the impact of hyperparameters and learn to tune them effectively.
- You are **not** allowed to use neural network libraries like PyTorch, Tensorflow and Keras.
- You are also **not** allowed to add, or remove any functions or classes, or even modify their names.
- You are also **not** allowed to change the signature (list of input attributes) of each function.
- Please note that this code may take several hours to run on one CPU. You will not find any gain by requesting GPUs on Mill.
- Finally, make sure that the gradients are structured according to their respective dimensions. Perform error handling if a dimensionality mismatch is observed.

Problem 1 Primitives

5 points

1. **LINEAR BASIS:** Implement a linear function in `LINEAR()` and its gradient in `LINEAR_GRAD()` (1 point)

- X is a $K \times 1$ vector, and W is a $M \times K$ matrix (Note that M is a hyperparameter).
- *Linear function:* $Y = W \cdot X$ is a $M \times 1$ vector.
- *Gradient of Linear function:* Return one of the following two gradients as needed:

$$\nabla_X Y = W \text{ (dimension } = M \times K), \quad \text{and} \quad \nabla_W Y = \begin{bmatrix} X & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & X \end{bmatrix}_{M \times M \times K}$$

2. **ACTIVATION FUNCTIONS:** Implement four activation functions, namely ReLU in `RELU()`, its gradient in `RELU_GRAD()`, logistic function in `LOGISTIC()` and its gradient in `LOGISTIC_GRAD()` (2 points)

Note: Let x_i be one of the entries in X . Then, activation functions are typically defined on each entry in X , i.e. $y_i = \sigma(x_i)$ for all $i = 1, \dots, M$

ReLU Activation:

- *ReLU function:* $y_i = \begin{cases} x_i, & \text{if } x_i \geq 0, \\ 0, & \text{otherwise.} \end{cases}$
- *Gradient of ReLU function:* $\nabla_{x_j} y_i = \begin{cases} 1, & \text{if } i = j \text{ and } y_i \geq 0, \\ 0, & \text{otherwise.} \end{cases}$

Note: The above definition includes the subgradient of ReLU at $x = 0$. Also, $\nabla_x y$ is $M \times M$ dimensional, and does not equate to identity matrix for every y .

Logistic Function:

- *Logistic function:* $y_i = \frac{1}{1 + \exp(-x_i)}$ for all $i = 1, \dots, M$.
- *Gradient of Logistic Function:* $\nabla_x y = \begin{bmatrix} -y_1(1-y_1) & 0 & \dots & 0 \\ 0 & -y_2(1-y_2) & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & -y_M(1-y_M) \end{bmatrix}_{M \times M}$

3. **LOSS FUNCTIONS:** In this primitive alone, assume that the true label Y , and therefore, the neural network output \hat{Y} are binary variables (i.e. $M = 1$). Return an error message if this condition does not hold true. Implement the binary cross-entropy loss in `CROSSENTROPY()`, and its gradient in `CROSSENTROPY_GRAD()`. (2 points)

- *Binary Cross Entropy:* $z = l(y, \hat{y}) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$
- *Gradient of Binary Cross Entropy:* $\nabla_{\hat{y}} z = \frac{1 - y}{1 - \hat{y}} - \frac{y}{\hat{y}}$ (a scalar quantity)

Note: Be aware of the $\frac{0}{0}$ case. A typical solution is to add a very small number $\delta = 10^{-6}$ to both numerator and denominator.

Problem 2 NN-2 Model

5 points

Using primitive functions defined in Problem 1, implement a two-layer NN in `NN2` class, with the following definitions:

1. **Forward pass:** Implement in `NN2.FORWARD()` definition. (1 point)

- Function: $\hat{y} = \sigma_2(\mathbf{w}_2^T \cdot \sigma_1(W_1 \cdot \mathbf{x}))$
- Assume $\sigma_2(\cdot)$ is a logistic function, and $\sigma_1(\cdot)$ a ReLU function.
- Assume W_1 is a $M \times K$ matrix, and \mathbf{w}_2 is a $M \times 1$ vector.

Backward Pass: Implement in `NN2.BACKWARD()` definition. (3 points)

- Let $z_1 = W_1 \cdot \mathbf{x}$, $\tilde{z}_1 = \sigma_1(z_1)$, and $z_2 = \mathbf{w}_2^T \cdot \tilde{z}_1$. Then, $\hat{y} = \sigma_2(z_2)$.

- Gradient Computation (Backpropagation):

$$* \quad \nabla_{W_1} \ell = \nabla_{\hat{y}} \ell \cdot \nabla_{z_2} \hat{y} \cdot \nabla_{\tilde{z}_1} z_2 \cdot \nabla_{z_1} \tilde{z}_1 \cdot \nabla_{W_1} z_1 = \left[\frac{\partial \ell}{\partial W_1(i, j)} \right] \in \mathbb{R}^{1 \times M \times K}$$

$$* \quad \nabla_{w_2} \ell = \nabla_{\hat{y}} \ell \cdot \nabla_{z_2} \hat{y} \cdot \nabla_{w_2} z_2 = \left[\frac{\partial \ell}{\partial w_2(m)} \right] \in \mathbb{R}^{1 \times M}$$

- Hint: You may use `NUMPY.DOT()` to compute the product of gradients.

Empirical loss gradient: Implement in `NN2.EMP_LOSS_GRAD()` definition. (1 point)

- Given a training data $(x_1, y_1), \dots, (x_N, y_N)$, the empirical risk is given by

$$L_N = \frac{1}{N} \sum_{i=1}^N \ell(y_i, \hat{y}_i).$$

- The gradient of empirical risk is given by

$$\nabla_w L_N = \frac{1}{N} \sum_{i=1}^N \nabla_w \ell(y_i, \hat{y}_i).$$

- **Note:** Everytime the optimization algorithm updates w , the gradient of loss function needs to be computed since \hat{y} changes accordingly.

Problem 3 Optimization Algorithms

5 points

- **GD:** Implement in `GRADIENT_DESCENT()` function.
 - Hyperparameter: Learning rate ϵ , Number of iterations R
 - Initialize $W^0 \sim \mathcal{N}(0, \sigma^2 I)$.
 - In the r^{th} iteration, compute the gradient of empirical loss with respect to W^{r-1} using `emp_loss_grad` function in the model class.
 - Compute the update step: $W^{(r)} = W^{(r-1)} - \epsilon \cdot \nabla L_N(W^{(r-1)})$
 - **Bonus Points:** You will receive one extra point for each one of the following optimization algorithms: SGD, AdaM

Problem 4 Classification on MNIST¹ Data

5 points

1. Data Preprocessing on MNIST:

(2 points)

¹Original Source: <http://yann.lecun.com/exdb/mnist/>

- MNIST data comprises of 70,000 images of handwritten digits from 0 to 9 (10 label classes), where each image has 28×28 pixels of gray-scale values ranging from 0 (black) to 1 (white). Flatten each image (28×28 matrix) into one 784×1 vector. Append a '-1' to the end of this array to get the 785×1 input, i.e. $K = 785$.
 - Convert these 10-ary labels into a binary label, where the outcome is '1' if the original image label is an **even** number, and '0' otherwise.
 - Partition the entire dataset into $T = 10,000$ test samples and the remaining as training samples.
2. **Training on MNIST:** Train NN-2 model on the training portion of the pre-processed MNIST dataset by choosing appropriate hyperparameters M , ϵ , R and the initial weights for W_1 and w_2 . (1 point)
- Note:** Your model performance depends on how well you choose your hyperparameters.
3. **Testing on MNIST:** Validate the performance of the trained NN-2 model using the testing portion of the pre-processed MNIST dataset. Report your performance in terms of accuracy, which is defined as

$$Acc = \frac{1}{|\text{Test Samples}|} \sum_{i \in \text{Test Samples}} \mathbb{1}(y_i = \hat{y}_i),$$

where $\mathbb{1}(A)$ is a indicator function that returns a value '1', when A is true. (2 points)

Note: An acceptable model produces at least 75% accuracy.